

# Taehoon Hwang

[taehoon.h@columbia.edu](mailto:taehoon.h@columbia.edu) | [linkedin.com/in/taehoon-hwang](https://linkedin.com/in/taehoon-hwang) | [taehoonhwang.net](http://taehoonhwang.net)

## EDUCATION

### Columbia University

B.S. in Computer Science; GPA: 3.98 / 4.33

New York, NY

Aug. 2025 – Present

**Relevant Courses** Computer Science Theory, Programming Languages and Translators,

Introduction to Databases, Projects in Computer Science, Theory ML Interaction,

Natural Language Processing

### Purdue University

B.S. in Computer Science, Minor in Mathematics; GPA: 4.00 / 4.00

West Lafayette, IN

Aug. 2023 – May. 2025

**Relevant Courses** Linear Algebra, Probability, Discrete Mathematics, Computer Architecture,

Data Structures, Statistical Methods, Competitive Programming I & II,

Foundations of Deep Learning, Data Mining & Machine Learning,

Introduction to Artificial Intelligence

## EXPERIENCE

### Research Assistant

Sep. 2025 – Present

Columbia University

New York, NY

- Working on fine-grained advantage computations for reinforcement learning to improve GRPO methodologies

### Artificial Intelligence Research Intern

May 2025 – Aug. 2025

Asteromorph

Seoul, South Korea

- Worked with fine-tuning and hosting LLMs (>200B parameters) on multi-GPU servers using vLLM and llama.cpp
- Developed an LLM fine-tuning framework by writing kernel code with PyTorch, Unslloth, and Hugging Face
- Deployed a parallel LLM fine-tuning scheduler and metric dashboard using Ray, Grafana, Prometheus, and Docker

### Teaching Assistant

Jan. 2025 – May. 2025

Purdue University

West Lafayette, IN

- Held office hours and lab sessions for students in Discrete Mathematics and Programming in C courses
- Reviewed and conducted problem solving sessions for discrete mathematics on first order logic, set theory, etc
- Answered students' questions during lab sessions on UNIX systems, file IO, dynamic memory allocation, etc

### Research Assistant

Jan. 2024 – Dec. 2024

Purdue University

West Lafayette, IN

- Engineered demand forecasting pipeline integrating probability curves with LSTM encoding for enhanced accuracy
- Optimized inference storage by achieving a reduction of >99% compared to traditional deep learning models
- Conducted latent space analysis and illustrated composite encoding of demand features using LSTM models

### Machine Learning Intern

Sep. 2023 – Aug. 2024

Quantum Research Sciences

West Lafayette, IN

- Employed ML algorithms and statistical techniques to deploy software for the United States Air Force
- Produced Monte Carlo simulations to test and tune quantum algorithms for inventory management
- Developed end-user interface for data analytics in quantum inventory management software

### Student Researcher

Mar. 2022 – Dec. 2022

Seoul Science High School

Seoul, South Korea

- Explored enhancements in real-time landmark tracking algorithms using machine and deep learning
- Improved Google MediaPipe performance by reducing untracked frames by 91.7% through U-Net segmentation
- Designed a versatile multi-headed software solution adaptable to various hardware specifications

### Research Assistant

May 2021 – Dec. 2021

Sungkyunkwan University

Seoul, South Korea

- Engaged with research team to boost open set recognition in ResNet models by 9% via latent space manipulation
- Developed and deployed a remote automated testing and optimization framework for models using PyTorch
- Presented research findings at the Korean Science High School R&E Conference

**Student Researcher** Apr. 2020 – Jun. 2021  
*Seoul Science High School* Seoul, South Korea

- Investigated topics on advertisement classification and recommendation using OCR and clustering algorithms
- Refined k-means clustering to enhance classification accuracy, adapting to the temporal dynamics of user interests
- Demonstrated algorithm adaptability to temporal features through rigorous testing and analysis

## PROJECTS

---

<b>Stack Rotation</b>	May 2025 – Aug. 2025
<ul style="list-style-type: none"><li>• Developed full-rank fine-tuning LoRA methodology with PyTorch using stochastically sampled rotation matrices</li><li>• Achieved SOTA-comparable performance when compared to LoRA variants such as DoRA, VeRA, and HiRA</li><li>• Constructed parallel fine-tuning orchestrator using Ray, Grafana, Prometheus, and Docker</li></ul>	
<b>Rotationally Equivariant Spatio-temporal Prediction</b>	Feb. 2025 – Apr. 2025
<ul style="list-style-type: none"><li>• Developed spatio-temporal predictive models with rotational equivariance to adapt to out-of-distribution inputs</li><li>• Constructed training infrastructure using PyTorch Lightning and distributed high-performance computing nodes</li><li>• Compared methods for equivariance using steerable wavelet filters, G-CNNs, and equivariant attention modules</li></ul>	
<b>Financial Anomaly Detection with Modified Benford's Law</b>	Jan. 2025 – Mar. 2025
<ul style="list-style-type: none"><li>• Suggested novel alteration of Benford's Law applicable to stock return data based on the Student-Lévy process</li><li>• Proved algorithm utility with real-life stock data and demonstrated 60% improvement over conventional methods</li><li>• Utilized new algorithm to detect financial anomalies in stock returns based on fitted location-scale t-distributions</li></ul>	
<b>WUMT: Wavelet U-Net Motion Transformer</b>	Jul. 2024 – Mar. 2025
<ul style="list-style-type: none"><li>• Investigated spatio-temporal encoding methods of videos using discrete wavelet transforms and NAFNet blocks</li><li>• Utilized 4D motion tensor computations and predictions with transformers for U-Net latent space operations</li><li>• Built training infrastructure with PyTorch Lightning and leveraged Docker and MLFlow for streamlined research</li></ul>	
<b>RL Wildfire Optimization</b>	May 2024 – Dec. 2024
<ul style="list-style-type: none"><li>• Utilized Convolutional DQN models to optimize firefighting efforts and evacuation routes with wildfires</li><li>• Developed complex wildfire simulation incorporating population density, terrain, weather, and crowd dynamics</li><li>• Theorized dynamic custom action and state space utilizing one-hot encoding and action masks</li></ul>	
<b>Contextual-Diffusion</b>	Feb. 2023 – Aug. 2023
<ul style="list-style-type: none"><li>• Improved cohesion and features of Stable Diffusion model output utilizing spacial context from LLMs</li><li>• Developed pipeline incorporating Mask-RCNN models as translation layers between LLMs and Stable Diffusion</li><li>• Generated and annotated image segmentation datasets for supervised learning of Mask-RCNN models</li></ul>	
<b>Minimax-based Animal Shogi AI</b>	Mar. 2021 – Nov. 2021
<ul style="list-style-type: none"><li>• Programmed an Animal Shogi bot using a minimax algorithm with alpha-beta pruning in C++</li><li>• Implemented interactable game GUI and cross-language translation mechanism with Python</li><li>• Demonstrated bot with live play-testing at school festival to 350+ students and faculty</li></ul>	

## ACTIVITIES

---

<b>Purdue Hackers Club Member</b>	Jan. 2025 – May 2025
<i>Purdue University</i>	West Lafayette, IN
<ul style="list-style-type: none"><li>• Participated in weekly 'hack' sessions and developed projects involving full-stack web development.</li><li>• Developed website for course scheduling at Purdue with Next.js, MongoDB, and Python among team of 3.</li></ul>	
<b>Purdue IGDC Club Member</b>	Jan. 2024 – May 2025
<i>Purdue University</i>	West Lafayette, IN
<ul style="list-style-type: none"><li>• Participated in various in-person game development sessions with other club members.</li><li>• Developed game with Unity centered around creating gravitational wells to enduce fast-paced gameplay.</li><li>• Gave feedback on other games being developed and participated in play testing.</li></ul>	

<b>ICPC 2023 ECNA Regional Participant</b>	Sep. 2023
<i>Purdue University</i>	<i>West Lafayette, IN</i>

- Ranked in the Top 20 in Purdue's ICPC participant selection contest
- Lead team of 3 in the ICPC East Central North America regional contest among 5 other Purdue ICPC teams
- Solved problems utilizing competitive programming techniques such as dynamic programming and backtracking

<b>Hello World Hackathon Participant</b>	Sep. 2023
<i>Purdue University</i>	<i>West Lafayette, IN</i>

- Led a team of 4 in a 24-hour hackathon driving project development and collaboration
- Built a software pipeline utilizing LLMs to deliver personalized dietary text feedback to users
- Built a full-stack web app with a React front-end and ExpressJS back-end integrating MongoDB RestAPIs

<b>ML@Purdue Club Member</b>	Mar. 2023 – May 2025
<i>Purdue University</i>	<i>West Lafayette, IN</i>

- Engaged in club discourse on AI and ML topics for research and software development.
- Participated in paper reading and workshops on various topics, such as transformers, GNNs, etc.

<b>Compute Infrastructure Manager</b>	Nov. 2021 – Feb. 2022
<i>Seoul Science High School</i>	<i>Seoul, South Korea</i>

- Allocated funding and purchased a multi-GPU cloud computing server for deep learning model training purposes.
- Managed and tested training infrastructure for research by setting up various software such as job managers.

<b>Published Author on Neural Network Fundamentals</b>	Sep. 2021 – Mar. 2022
<i>Barun Books Co.,Ltd</i>	<i>Seoul, South Korea</i>

- Authored a 160-page book on neural networks covering backpropagation, gradient descent, and parallelization
- Distributed to 13 retailers, selling 300+ copies in the first year
- Ranked as a top 4 entry in the "Weekly Top Releases" by the second-largest South Korean book retailer

<b>Artificial Intelligence Lecturer</b>	May 2021 – Apr. 2023
<i>Seoul Science High School</i>	<i>Seoul, South Korea</i>

- Presented 8 deep learning-focused lectures to an audience of 200+ students and faculty
- Delivered lectures on topics including linear algebra, conventional neural networks, and transformers
- Conducted 10 additional interdisciplinary volunteer lectures focused on mathematics and artificial intelligence

<b>Artificial Intelligence Club President</b>	Mar. 2021 – Feb. 2023
<i>Seoul Science High School</i>	<i>Seoul, South Korea</i>

- Authored 4 different entry exams on mathematical deep learning and essays on model selection
- Curated literature and project content with paper readings and specialized datasets
- Supervised and managed the acquisition of a multi-GPU deep learning server for club research and projects

## HONORS & AWARDS

---

### Dean's List

*Columbia University, 2025*

### Purdue Office of Undergraduate Research Grant

*Purdue University, 2024* \$ 500

### Dean's List & Semester Honors

*Purdue University, 2023 & 2024 & 2025*

### Excellence Award in the Gifted School Research Conference

*Pohang University of Science & Technology, 2022*

### Top Award for R&E Research in Computer Science

*Seoul Science High School, 2022*

### Azure AI Fundamentals Certification

*Microsoft, 2022*

### Top Award for Independent Research on Computer Science

*Seoul Science High School, 2021*

## TECHNICAL SKILLS

---

**Languages & Frameworks:** C/C++, Express.js, Flask, HTML/CSS, Java, JavaScript, Python, React, x86 Assembly

**Dev & Research Tools:** Docker, Git/GitHub, IntelliJ, LaTeX, MLflow, Overleaf, PyCharm, Slurm, VS Code

**Libraries:** Keras, NumPy, OpenCV, Pandas, PyTorch, Scikit-learn, TensorFlow, Transformers, Unislot, trl, vLLM